

Welcome to Machine Learning Module



Huntsman Faculty & Staff Directory

Economics and Finance Department

Management Department

Management Information Systems
Department

Marketing and Strategy Department

School of Accountancy

Pedram Jahangiry

Visiting Assistant Professor

Department(s):
Economics and Finance Department

Educational Background

PhD, Economics, Arizona State University, 2017
Master, Economics, Simon Fraser University, 2013
MBA, Sharif University, 2012
Industrial Engineering, IUST, 2009

Biography

Pedram Jahangiry, PhD is an assistant professor in the Economics and Finance Department of the Jon M. Huntsman School of Business at Utah State University. He recently passed level III in the CFA program. Prior to joining the Huntsman School in 2018, Pedram was a research associate within Financial Modeling Group at BlackRock NYC. His research is involved in machine learning applications in finance, empirical asset pricing, and factor models.

Show more




Contact Information


📍 Office Location:
[Eccles Business Building 507](#)
📞 Phone: [435.797.2345](tel:435.797.2345)
✉ Email: pedram.jahangiry@usu.edu

[Personal Website](#) [Curriculum Vitae](#)

- What is Machine Learning?
- Different major types
- Applications
- Key terminology
- Applied ML



Fancy plot: monthly stock returns



What do we learn?


Useful plot: market model coefficients

Fit a line between stock returns R_t and market returns M_t (SP).

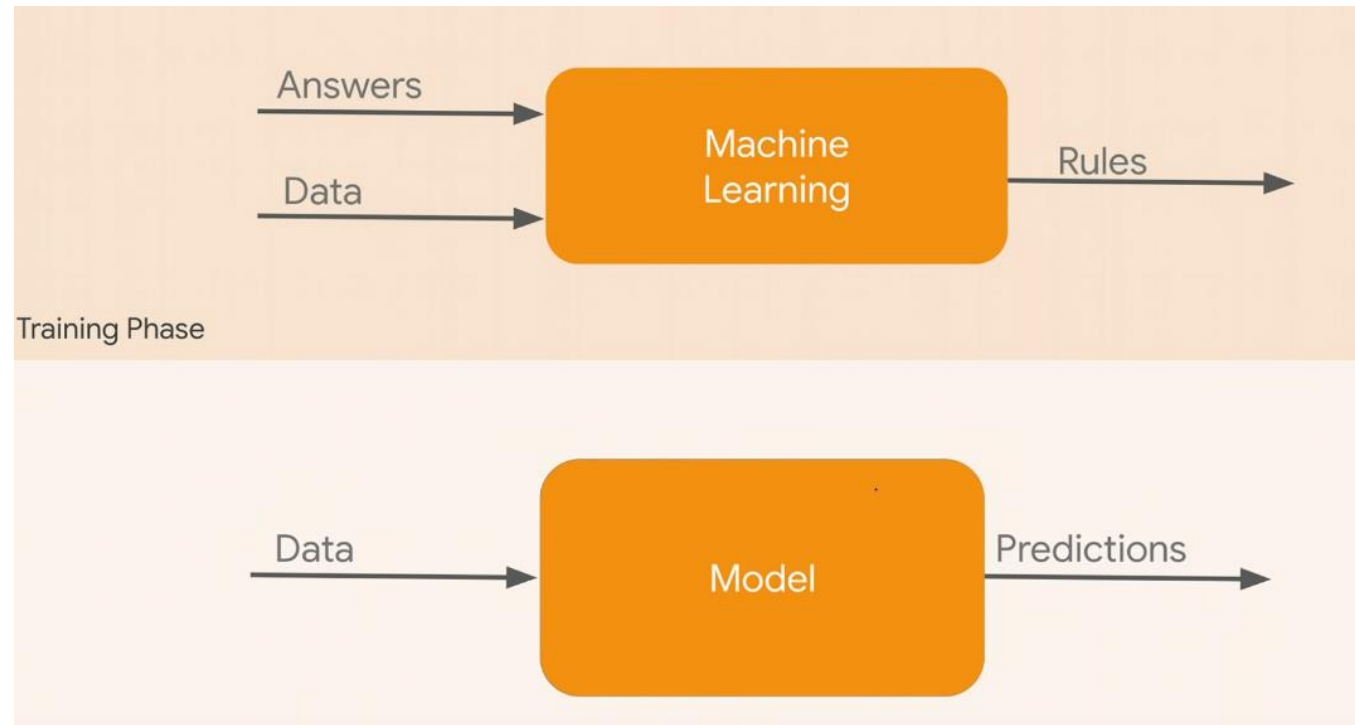
$$R_t \approx \alpha + \beta M_t$$

α is money you make regardless of what the market does.


β is the asset's sensitivity to broad market movements.



Machine Learning vs. Traditional Programming



What is Machine Learning?



What is Machine Learning?

Machine Learning is the field that thinks about how to **automatically** build robust **predictions** from **complex data**.


- ML is closely related to modern statistics (many ideas come from statisticians!)
- Statistics focus on **model inference**, ML focus on **predictive performance**
- The entire field of ML is calibrated against “**out-of-sample**” experiments that evaluate how well a model **trained** on one dataset (train data) will **predict** new data (test data).
- While there is a recent push to build more transparency into machine learning, wise practitioners avoid assigning **structural** meaning to the parameters of their fitted models.
- ML models are **black boxes** whose purpose is to do a good job in **predicting** a future that follows the **same patterns** as in past data.

Why you need to learn ML?

- To project the information in your data into a **low-dimensional space** that contains key insights for the decisions you need to make
- To give you the tools to quickly **turn messy data into useful information** that has direct relevance to business policy
- Big data + Machine learning = Data science
- Big data + Machine learning + Economics/Econometrics= **Business data science**

Understanding the major types of machine learning

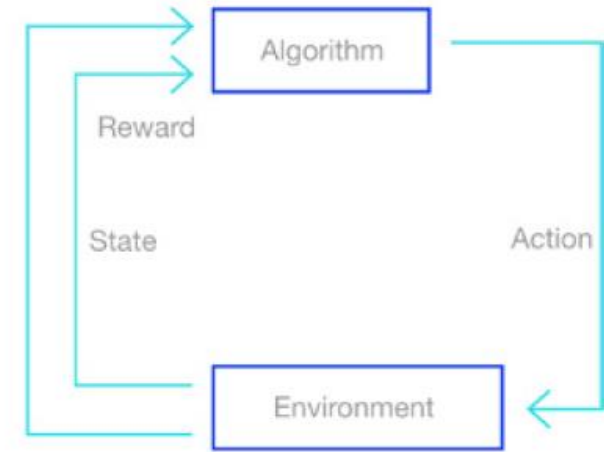
Supervised learning



Unsupervised learning



Reinforcement learning




What it is

An algorithm uses training data and feedback from humans to learn the relationship of given inputs to a given output (eg, how the inputs “time of year” and “interest rates” predict housing prices)

An algorithm explores input data without being given an explicit output variable (eg, explores customer demographic data to identify patterns)

An algorithm learns to perform a task simply by trying to maximize rewards it receives for its actions (eg, maximizes points it receives for increasing returns of an investment portfolio)

Machine Learning Branches



Regression

Regression models (both linear and non-linear) are used for predicting a real value, like salary for example. If your independent variable is time, then you are forecasting future values, otherwise – your model is predicting present but unknown values. Regression techniques vary from MLR to SVR and Boosted Trees.

5 Applications Of Regression in Real world



Future Prediction




To predict events yet to occur. For example predicting future sales, units consumed by consumer.

Process Optimization



To optimize business processes like a manager in call center would like to know how much call waiting time affects the number of complaints.

Market Forecast



Here Multivariate linear technique is used to forecast sales volumes, or market movement.

Human Resources



Linear regression is also used to predict the demographics and types of future work force in large companies. This way companies can create better hiring plans.

Study of Environment



Linear Regression has a wide applications in Environment Science. For example in Canada, statistical analysis is used to measure effect of oil spills on aquatic ecosystem.


Machine Learning Branches

Unlike regression where you predict a continuous number, you use classification to predict a category. There is a wide variety of classification applications from medicine to marketing. Classification models include linear models like Logistic Regression, SVM, and nonlinear ones like K-NN, Kernel SVM and Random Forests.



Classification

Machine Learning Branches



Clustering

Clustering is similar to classification, but the basis is different – in Clustering you don't know what you are looking for. When you use clustering algorithms on your dataset, unexpected things can suddenly pop up – like structures, clusters and groupings you would have never thought of otherwise.

5 Applications Of Clustering In Real World

Segmenting Customers



Used in marketing, where customers are subdivided into segments with each segment consisting of customers with similar characteristics.

Creating Newsfeeds



To cluster articles from various sources by their similarity – it can separate documents into disjoint clusters and hence helps decision making.

Pattern Recognition



This feature is readily used in medical field these days where algorithm automatically detect for segmentation of blood cells for leukemia detection or for cancer cells detection.

Cloud Storage



To enhance performance, capacity, or reliability, clustering distributes work loads to each server, and provides access to all files from any server regardless of the physical location of the file.

Recommender System



To recommend similar items for customers. To find similar items, comparison b/w the set of users who like each item is done, if a similar set of users like two different items, then the items themselves are probably similar!

Machine Learning Branches

Reinforcement learning algorithms include techniques like Thompson Sampling, Upper Confidence Bound and Q-Learning. These are used a lot when training machines to perform tasks such as walking. Desired outcomes provide the AI with reward, undesired – with punishment. Machines learn through trial and error.




Reinforcement Learning

Machine Learning Branches



Natural Language Processing

Teaching machines to understand what is said in spoken and written word is the focus of Natural Language Processing. Whenever you dictate something into your iPhone / Android device and it's converted to text – that's an NLP algorithm in action. Methods include decision trees, Markov processes, and more.




- Familiarize you with ML terminology
- Prepare you for job interviews
- Provide cheat sheets and resources
- Help you to hedge yourself against the next recession



Hedge yourself against
the next recession!

When Will Next U.S. Recession Begin?



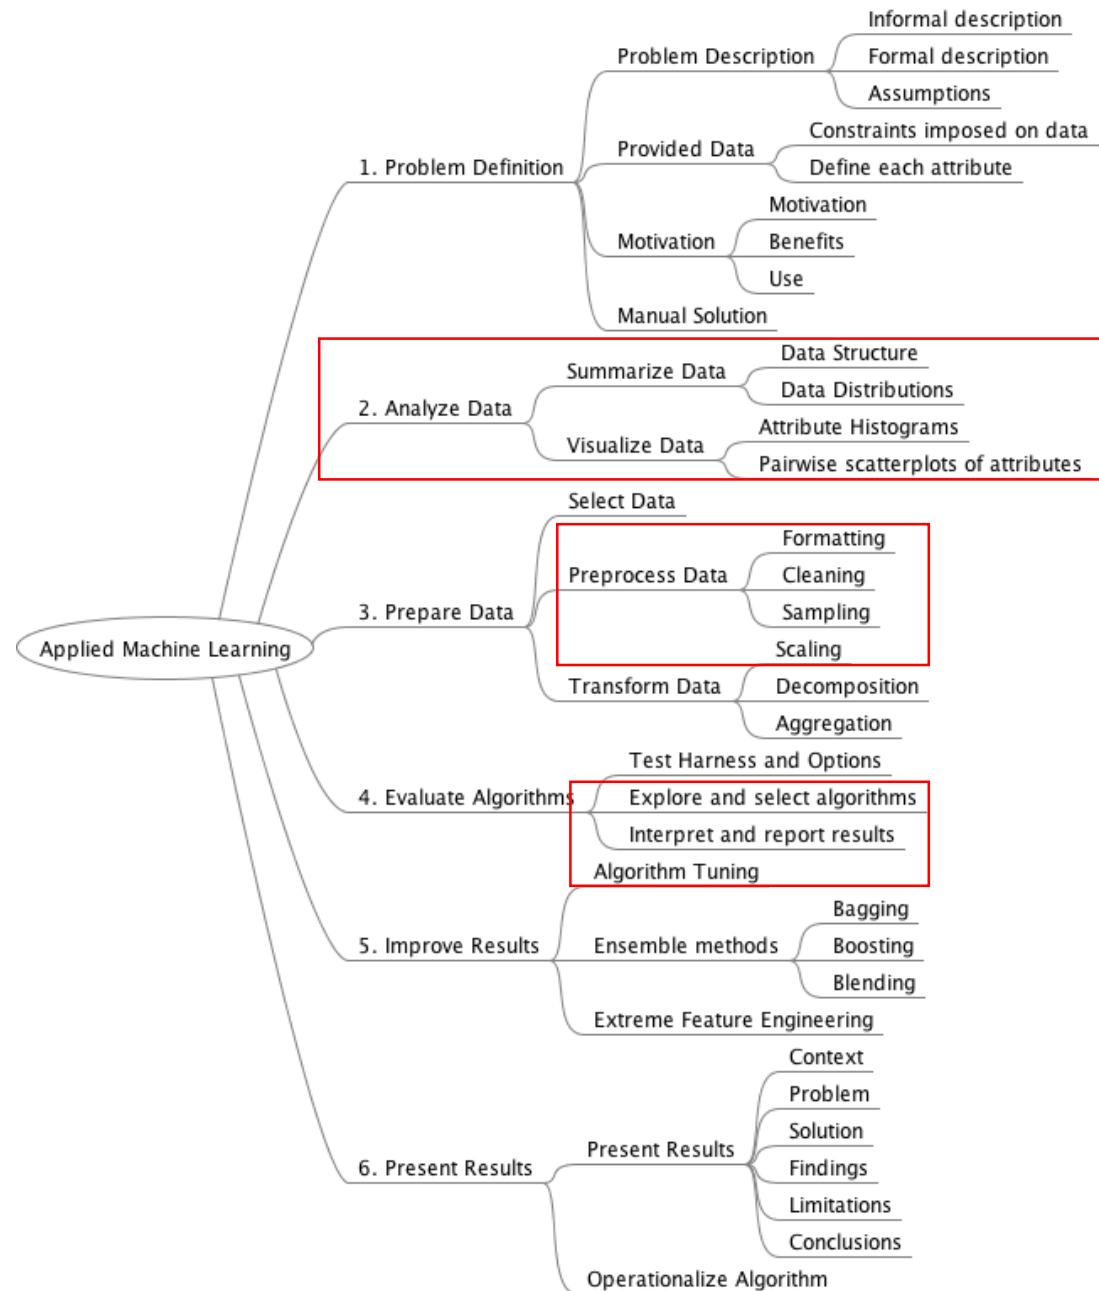
INDUSTRY

Machine Learning

Accommodation & Food	Agriculture	Banking & Insurance
Biotechnological & Life Sciences	Construction & Engineering	Education & Research
Emergency & Relief	Finance	Manufacturing
Government and Public Works	Healthcare	Media & Publishing
Justice, Law and Regulations	Miscellaneous	Accounting
Real Estate, Rental & Leasing	Utilities	Wholesale & Retail

- Accommodation & Food
 - Food
 - Restaurant
 - Accommodation
- Accounting
 - Machine Learning
 - Analytics
 - Textual Analysis
 - Data
 - Research and Articles
 - Websites
 - Courses
- Agriculture
 - Economics
 - Development
- Banking & Insurance
 - Consumer Financial
 - Management and Operations
 - Valuation
 - Fraud
 - Insurance and Risk
 - Physical
- Biotechnological & Life Sciences
 - General
 - Sequencing
 - Chemoinformatics and drug discovery
 - Genomics
 - Life-sciences
- Construction & Engineering
 - Construction
 - Engineering
 - Material Science
- Economics
 - General
 - Machine Learning
 - Computational
- Education & Research
 - Student
 - School
- Emergency & Relief
 - Preventative and Reactive
 - Crime
 - Ambulance
 - Disaster Management
- Finance
 - Trading & Investment
 - Data
- Healthcare
 - General
- Justice, Law and Regulations
 - Tools
 - Policy and Regulatory
 - Judicial
- Manufacturing
 - General
 - Maintenance
 - Failure
 - Quality
- Media & Publishing
 - Marketing
- Miscellaneous
 - Art
 - Tourism
- Physics
 - General
 - Machine Learning
- Physics
 - General
 - Machine Learning
- Government and Public Works
 - Social Policies
 - Election Analysis
 - Disaster Management
 - Politics
 - Charities
- Real Estate, Rental & Leasing
 - Real Estate
 - Rental & Leasing
- Utilities
 - Electricity
 - Coal, Oil & Gas
 - Water & Pollution
 - Transportation
- Wholesale & Retail
 - Wholesale
 - Retail


What we learn
in this class?



Classes and events



Why Python?



TOP PYTHON / SCALA / R LIBRARIES IN DATA SCIENCE

PYTHON	COMMITTS	CONTRIBUTORS	SCALA	COMMITTS	CONTRIBUTORS	R	COMMITTS	CONTRIBUTORS
TensorFlow	34 689	1 505	DEEPLARNING4J	22 938	223	H2O	22 947	104
dlib	22 811	1 099	Spark MLlib & ML	22 073	1 241	mlr	4 052	59
PYTORCH	11 763	668	PredictionIO	4 383	125	dmlc XGBoost	3 296	286
Keras	4 591	688	bigDL	2 358	50	caret	1 741	65
XGBoost	3 296	286	summingbird	1 780	33	LightGBM	1 125	82
CatBoost	1 746	63	DeepLearning.scala	1 874	15	Prophet	284	34
LightGBM	1 125	82	Spark + H2O SPARKLING WATER	1 416	33	quartz	4 070	152
dist-keras	1 125	5	Smile	1 195	30	d3 plotly	3 185	24
elephas	170	13	Conjecture	155	8	ggvis	2 159	21
spark-deep-learning	68	11	VEGAS	215	16	DT DataTables	1 921	22
elis	929	6	Breeze-viz	29	3	rCharts	638	11
matplotlib	25 984	727	akka	22 298	552	Corrplot	335	10
Bokeh	17 066	298	kafka	5 145	442	lattice	781	3
d3 plotly	2 910	47	Breeze	3 431	89	dplyr	4 990	151
Seaborn	2 093	85	Epic	1 790	15	data.table	3 416	50
pydot	169	12	Puck	536	1	lubridate	1 501	51
SciPy	19 255	613	saddle	184	10	jsonlite	914	12
NumPy	18 162	654	Scalalab	44	1	kern	5 528	105
Natural Language Toolkit	13 053	238	Slick	2 037	107	redbox	2 373	64
Statsmodels Statistics in Python	10 244	153	ADAM	1 775	64	slidyfy	302	6
spaCy	8 655	226						
gensim	3 610	275						
pandas	17 243	1 179						
Scrapy	6 643	285						

LEGEND: Machine Learning Visualization Mathematics & Engineering Data Manipulation & Analysis Reproducible Research

Introduce your self!

What project are you working on?

