

Regression

September 16, 2019

0.1 Pedram Jahangiry, Fall 2019

1 Regression Analysis:

A linear Regression is a **linear approximation** of a **causal relationship** between two or more variables

1.1 Multiple Regression

First we need to import the libraries:

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

sns.set() #if you want to use seaborn themes with matplotlib functions
```

1.1.1 Data Preprocessing

```
[2]: df = pd.read_csv("wage.csv")
df.head()
```

```
[2]:
```

	wage	hours	IQ	educ	exper	tenure	age	married	black	meduc	\
0	769000.0	40	93	12	11	2	31	1	0	8.0	
1	808000.0	50	119	18	11	16	37	1	0	14.0	
2	825000.0	40	108	14	11	9	33	1	0	14.0	
3	650000.0	40	96	12	13	7	32	1	0	12.0	
4	562000.0	40	74	11	14	5	34	1	0	6.0	

	feduc
0	8.0
1	14.0
2	14.0
3	12.0
4	11.0

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 935 entries, 0 to 934
Data columns (total 11 columns):
wage      935 non-null float64
hours     935 non-null int64
IQ        935 non-null int64
educ      935 non-null int64
exper     935 non-null int64
tenure    935 non-null int64
age       935 non-null int64
married   935 non-null int64
black     935 non-null int64
meduc     857 non-null float64
feduc     741 non-null float64
dtypes: float64(3), int64(8)
memory usage: 80.4 KB
```

```
[4]: df.describe().T
```

```
[4]:
```

	count	mean	std	min	25%	50%	\
wage	935.0	957945.454545	404360.822474	115000.0	669000.0	905000.0	
hours	935.0	43.929412	7.224256	20.0	40.0	40.0	
IQ	935.0	101.282353	15.052636	50.0	92.0	102.0	
educ	935.0	13.468449	2.196654	9.0	12.0	12.0	
exper	935.0	11.563636	4.374586	1.0	8.0	11.0	
tenure	935.0	7.234225	5.075206	0.0	3.0	7.0	
age	935.0	33.080214	3.107803	28.0	30.0	33.0	
married	935.0	0.893048	0.309217	0.0	1.0	1.0	
black	935.0	0.128342	0.334650	0.0	0.0	0.0	
meduc	857.0	10.682614	2.849756	0.0	8.0	12.0	
feduc	741.0	10.217274	3.300700	0.0	8.0	10.0	

	75%	max
wage	1160000.0	3078000.0
hours	48.0	80.0
IQ	112.0	145.0
educ	16.0	18.0
exper	15.0	23.0
tenure	11.0	22.0
age	36.0	38.0
married	1.0	1.0
black	0.0	1.0
meduc	12.0	18.0
feduc	12.0	18.0

```
[5]: df.isna().sum()
# Alternatively we could use isnull() from pandas.
# pd.isnull(df).sum()
```

```
[5]: wage          0
hours            0
IQ              0
educ            0
exper           0
tenure          0
age             0
married         0
black           0
meduc          78
feduc         194
dtype: int64
```

1. Because the number of NAs in feduc and meduc is greater than 5% of the observations, we should not keep them in the regression
2. For practice purposes, I will keep the meduc in the features.

```
[6]: df.drop('feduc', axis=1, inplace=True) #why do we need inplace?
```

```
[7]: # we will replace the missing meduc with median. Because the data is left-
    →skewed
# mean is not a good representation of the central tendency measure.


df['meduc'].fillna(df['meduc'].median(),axis=0, inplace=True )
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 935 entries, 0 to 934
Data columns (total 10 columns):
wage          935 non-null float64
hours         935 non-null int64
IQ            935 non-null int64
educ          935 non-null int64
exper         935 non-null int64
tenure        935 non-null int64
age           935 non-null int64
married       935 non-null int64
black         935 non-null int64
meduc         935 non-null float64
dtypes: float64(2), int64(8)
memory usage: 73.1 KB
```

1.1.2 Data visualization


```
[8]: plt.figure(figsize=(12,5))  
sns.distplot(df['wage'], bins=30 , rug=True)
```

[8]: <matplotlib.axes._subplots.AxesSubplot at 0x26f9d9bcc18>



```
[9]: # Need to do log transformation to avoid potential heteroskedasticity  
df['lwage'] = np.log(df['wage'])  
sns.distplot(df['lwage'], bins=30 , rug=True)
```


[9]: <matplotlib.axes._subplots.AxesSubplot at 0x26f9dd36a90>



```
[10]: df.drop('wage', axis=1, inplace=True)
```


```
[11]: plt.figure(figsize=(12,5))
sns.distplot(df['IQ'], bins=30 ,color='r' , rug=True)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x26f9de1a780>
```



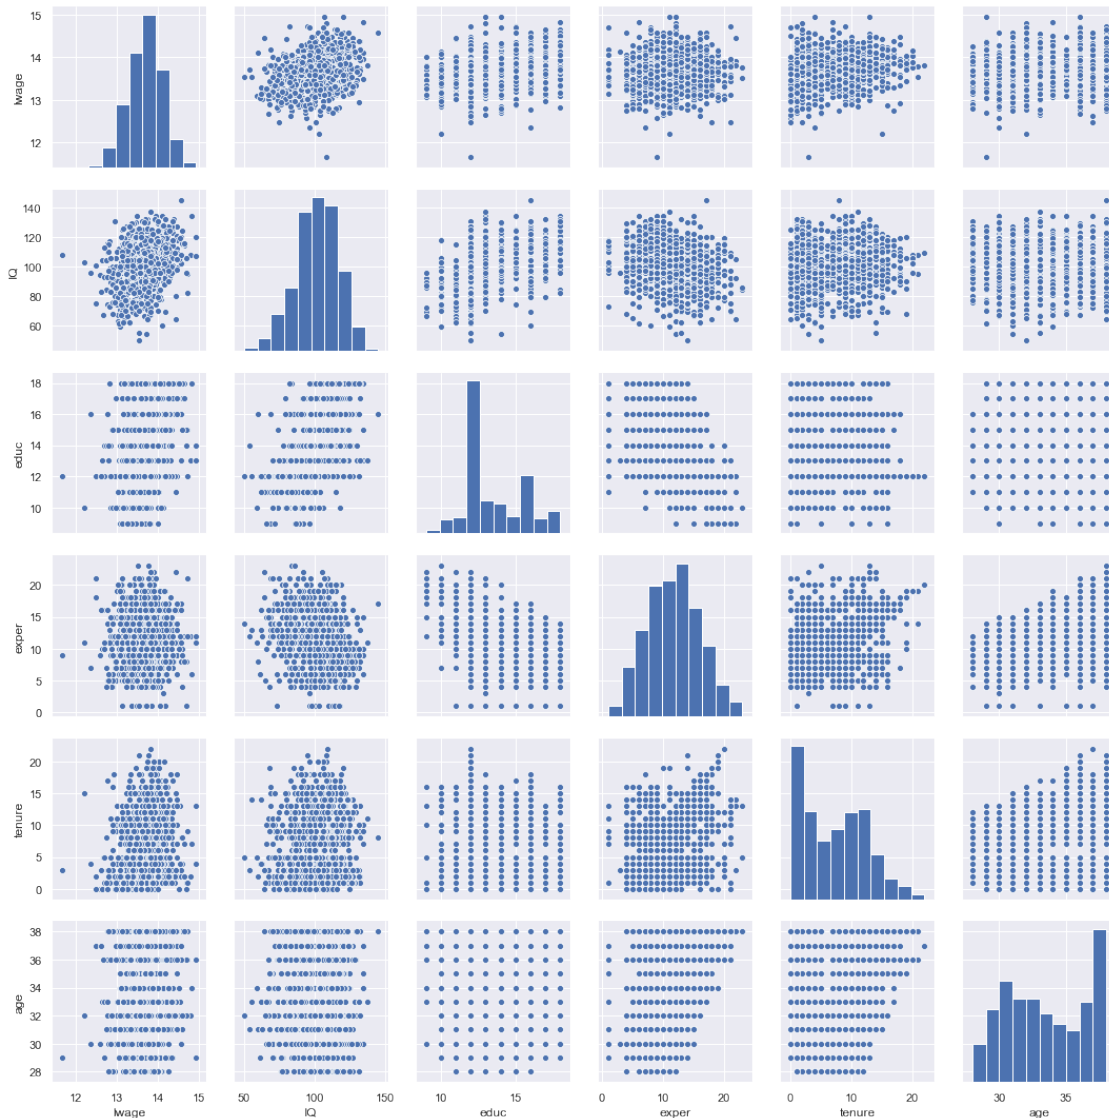
```
[12]: plt.figure(figsize=(12,5))
sns.heatmap(df.corr(), cmap='coolwarm',annot=True)
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x26f9e132a58>
```



```
[13]: sns.pairplot(df[['lwage', 'IQ', 'educ','exper','tenure','age']])
```

[13]: <seaborn.axisgrid.PairGrid at 0x26f9e143940>



1.1.3 Dealing with dummy variables

```
[14]: df2 = df.copy() # make sure you put df.copy(). Why?
df2['status']=df2['married'].map({1:'Married', 0:'Single'})
df2.drop('married',axis=1 ,inplace=True)
df2.head()
```

```
[14]:   hours  IQ  educ  exper  tenure  age  black  meduc      lwage  status
0     40   93   12     11        2   31      0      8.0  13.552846  Married
1     50  119   18     11       16   37      0     14.0  13.602317  Married
2     40  108   14     11        9   33      0     14.0  13.623139  Married
```

3	40	96	12	13	7	32	0	12.0	13.384728	Married
4	40	74	11	14	5	34	0	6.0	13.239257	Married

```
[15]: df2 = pd.get_dummies(df2, drop_first=True)
df2.head()
```

```
[15]:   hours  IQ  educ  exper  tenure  age  black  meduc  lwage  \
0     40  93   12    11      2   31     0    8.0  13.552846
1     50 119   18    11     16   37     0   14.0  13.602317
2     40 108   14    11      9   33     0   14.0  13.623139
3     40  96   12    13      7   32     0   12.0  13.384728
4     40  74   11    14      5   34     0    6.0  13.239257
```

```
   status_Single
0              0
1              0
2              0
3              0
4              0
```

in this example, since black and married are already dummy variables, we don't need to do any thing else.

1.1.4 Defining the variables and splitting the data

```
[16]: y = df['lwage']
X = df.drop('lwage', axis=1) # becareful inplace= False

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
→random_state=100)

len(X_train)/len(X)
```

```
[16]: 0.8
```

1.2 Linear Regression with StatsModels

Statsmodel is great for learning the theory of regression models. Also Statsmodel works perfectly with pandas dataframe. However, sklearn is a more **practical** package preferred by ML practitioners to apply regression analysis.

```
[17]: # Add a constant
X_test_wc = sm.add_constant(X_test)
X_train_wc = sm.add_constant(X_train)
```

```
C:\Users\Pedram\Anaconda3\lib\site-packages\numpy\core\fromnumeric.py:2389:
FutureWarning: Method .ptp is deprecated and will be removed in a future
version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
```

```
[18]: X_train_wc.head()
```

```
[18]:      const  hours   IQ  educ  exper  tenure  age  married  black  meduc
457      1.0    50  130   14    15      1   33         1     0   12.0
807      1.0    40  104   12    10      2   29         1     0   12.0
859      1.0    60  105   16    12      1   35         1     0   14.0
174      1.0    60  116   12     9      7   30         1     0   12.0
417      1.0    58  113   16     9      0   30         1     0   12.0
```

```
[19]: # Fit the model
model = sm.OLS(y_train,X_train_wc)
statsmodels_reg= model.fit()
```

```
[20]: statsmodels_reg.summary()
```

```
[20]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                        OLS Regression Results
=====
Dep. Variable:          lwage      R-squared:          0.228
Model:                  OLS      Adj. R-squared:       0.218
Method:                 Least Squares      F-statistic:      24.19
Date:                  Mon, 16 Sep 2019      Prob (F-statistic):  1.82e-36
Time:                  18:18:27      Log-Likelihood:     -296.16
No. Observations:      748      AIC:              612.3
Df Residuals:          738      BIC:              658.5
Df Model:               9
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                12.0552      0.200      60.323      0.000      11.663      12.448
hours                -0.0051      0.002     -2.771      0.006      -0.009     -0.001
IQ                   0.0038      0.001      3.460      0.001       0.002       0.006
educ                 0.0450      0.008      5.586      0.000       0.029       0.061
exper                0.0071      0.004      1.702      0.089      -0.001       0.015
tenure              0.0088      0.003      3.198      0.001       0.003       0.014
age                 0.0146      0.005      2.753      0.006       0.004       0.025
married             0.1633      0.043      3.773      0.000       0.078       0.248
black              -0.1707      0.043     -3.983      0.000      -0.255     -0.087
meduc               0.0109      0.005      2.085      0.037       0.001       0.021
=====
Omnibus:              13.666      Durbin-Watson:       2.227
Prob(Omnibus):         0.001      Jarque-Bera (JB):    18.162
Skew:                 -0.197      Prob(JB):            0.000114
Kurtosis:              3.654      Cond. No.            1.79e+03
=====
```

```
Warnings:
```


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 1.79e+03. This might indicate that there are strong multicollinearity or other numerical problems.
 """

Interpreting the results

```
[21]: # try
# statsmodels_reg.
statsmodels_reg.conf_int(alpha=0.01) # if you want to be more conservative set
→ alpha=0.01
```

```
[21]:
```

	0	1
const	11.539143	12.571348
hours	-0.009828	-0.000345
IQ	0.000962	0.006618
educ	0.024192	0.065791
exper	-0.003675	0.017885
tenure	0.001694	0.015917
age	0.000903	0.028227
married	0.051533	0.275097
black	-0.281414	-0.060039
meduc	-0.002603	0.024446

Graph of Actual vs. Predicted values

```
[22]: corr = round(y_train.corr(statsmodels_reg.fittedvalues), 2)
sns.scatterplot(x=y_train, y=statsmodels_reg.fittedvalues, alpha=0.6)
sns.lineplot(y_train, y_train)



plt.xlabel('Actual log wage', fontsize=14)
plt.ylabel('Prediced log wage', fontsize=14)
plt.title(f'Actual vs Predicted log wage:(Corr {corr})', fontsize=17)

plt.show()

sns.scatterplot(x=np.e**y_train, y=np.e**statsmodels_reg.fittedvalues, alpha=0.
→6)
sns.lineplot(np.e**y_train, np.e**y_train)


plt.xlabel('Actual wage', fontsize=14)
plt.ylabel('Prediced wage', fontsize=14)
plt.title(f'Actual vs Predicted wage:(Corr {corr})', fontsize=17)

plt.show()
```




Residuals vs Predicted values

```
[23]: sns.scatterplot(x=statsmodels_reg.fittedvalues, y=statsmodels_reg.resid ,  
    ↪alpha=0.6)  
  
plt.xlabel('Predicted log wage', fontsize=14)  
plt.ylabel('Residuals', fontsize=14)  
plt.title('Residuals vs Fitted Values', fontsize=17)  
  
plt.show()
```



```
[24]: resid_mean = round(statsmodels_reg.resid.mean(), 3)  
resid_skew = round(statsmodels_reg.resid.skew(), 3)  
  
sns.distplot(statsmodels_reg.resid)  
plt.title(f'Log price model: residuals Skew ({resid_skew}) Mean ({resid_mean})')  
plt.show()
```



```
[25]: # Mean Squared Error
MSE = round(statsmodels_reg.mse_resid, 3)
MSE
```

```
[25]: 0.131
```

1.3 Linear Regression with Scikit-Learn

So far we have only worked with data frames using Pandas. Now we may need to transform our data into arrays by using numpy because sklearn uses arrays instead of data frames.

Scikit-learn is a very powerful package enabling you to do almost everything in machine learning including Regression, Classification, Clustering, SVM and Dimensionality reduction. However, I don't recommend sklearn for deep learning algorithms. Pytorch, Tensorflow and Keras are better alternatives for deep learning.

Note: with sklearn, we don't need to add constants manually.

```
[26]: from sklearn.linear_model import LinearRegression
```

```
[27]: sklearn_reg = LinearRegression()
```

```
[28]: sklearn_reg.fit(X_train, y_train)
```

```
[28]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

1.3.1 Generating results

try reg.

```
[29]: # The coefficients of the regression
sklearn_reg.coef_
```

```
[29]: array([-0.00508645,  0.00379004,  0.04499157,  0.00710532,  0.00880574,
          0.01456483,  0.16331493, -0.17072619,  0.01092151])
```

```
[30]: # The intercept of the regression
sklearn_reg.intercept_
```

```
[30]: 12.055245334454883
```

```
[31]: # The R-squared of the regression
sklearn_reg.score(X_train,y_train)
```

```
[31]: 0.2278261765190378
```

```
[32]: X_train.head()
```

```
[32]:      hours  IQ  educ  exper  tenure  age  married  black  meduc
457      50  130   14    15        1   33         1     0   12.0
807      40  104   12    10        2   29         1     0   12.0
859      60  105   16    12        1   35         1     0   14.0
174      60  116   12     9        7   30         1     0   12.0
417      58  113   16     9        0   30         1     0   12.0
```

```
[33]: # If we want to find the Adjusted R-squared we can do so by knowing the R2, the
      ↪ # observations, the # features
R2 = sklearn_reg.score(X_train,y_train)
n = X_train.shape[0]
p = X_train.shape[1]
```

```
# We find the Adjusted R-squared using the formula
adjusted_R2 = 1-(1-R2)*(n-1)/(n-p-1)
adjusted_R2
```

```
[33]: 0.21840942257414797
```

```
[34]: # Let's create a new data frame with the names of the features
reg_summary = pd.DataFrame(data = X_train.columns.values, columns=['Features'])
reg_summary ['Coefficients'] = sklearn_reg.coef_
reg_summary
```

```
[34]:   Features  Coefficients
0    hours    -0.005086
1      IQ     0.003790
2    educ     0.044992
3    exper     0.007105
4   tenure     0.008806
5     age     0.014565
6  married     0.163315
```

```
7    black    -0.170726
8    meduc     0.010922
```

1.3.2 Further Diagnostic tests

Multicollinearity sklearn does not have a built-in way to check for multicollinearity. The main reason is that this is an issue well covered in statistical frameworks and not in ML ones. However, we can use statsmodels to run the VIF test.

```
[35]: X_train.columns.values
```

```
[35]: array(['hours', 'IQ', 'educ', 'exper', 'tenure', 'age', 'married',
          'black', 'meduc'], dtype=object)
```

```
[36]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
[37]: collinearity = X_train[['hours', 'IQ', 'educ', 'exper', 'tenure', 'age', 'married',
          → 'black', 'meduc']]
VIF = pd.DataFrame()
```

```
[38]: VIF["Features"] = collinearity.columns
VIF["VIF"] = [variance_inflation_factor(collinearity.values, i) for i in
          → range(collinearity.shape[1])]

VIF
```

```
[38]:   Features      VIF
0    hours  33.859663
1      IQ  60.857443
2    educ  68.139961
3   exper  15.038459
4  tenure   3.456928
5    age  116.108410
6 married   9.420645
7   black   1.273572
8   meduc  19.384284
```

heteroskedasticity

```
[39]: # using the White test for example:
from statsmodels.stats.diagnostic import het_white
```

```
[40]: hetero = X_train_wc
F_stat_pvalue = het_white(statsmodels_reg.resid.values, hetero.values,
          → retres=False)[3]
F_stat_pvalue
```

```
[40]: 0.0014769435236610107
```

1.3.3 Testing

Once we have trained and fine-tuned our model, we can proceed to testing it. Testing is done on a dataset that the algorithm has never seen

```
[41]: # Using our statmodels results.  
statsmodels_reg.predict(X_test_wc).head(4)
```

```
[41]: 143    13.725665  
      229    13.753259  
      116    13.664840  
      134    13.565368  
      dtype: float64
```

```
[42]: # Using our sklearn reg.  
sklearn_reg.predict(X_test)[0:4]
```

```
[42]: array([13.72566539, 13.75325897, 13.66484037, 13.56536839])
```

```
[43]: y_hat_test = sklearn_reg.predict(X_test)  
log_predictions = pd.DataFrame( {'Actuals':y_test , 'Predictions': y_hat_test})  
predictions = np.exp(log_predictions)  
predictions.tail()  
  
# You can reset the index if you wish. How?
```

```
[43]:
```

	Actuals	Predictions
471	511000.0	7.827895e+05
191	840000.0	1.093139e+06
688	865000.0	8.161840e+05
10	930000.0	1.435088e+06
420	1850000.0	8.345331e+05


```
[44]: # Additionally, we can calculate the difference and percentage difference  
      ↳between the targets and the predictions  
predictions['Residuals'] = predictions['Predictions'] - predictions['Actuals']  
predictions['Difference%'] = np.absolute(predictions['Residuals']/  
      ↳predictions['Actuals']*100)  
predictions.round().tail(5)
```

```
[44]:
```


	Actuals	Predictions	Residuals	Difference%
471	511000.0	782789.0	271789.0	53.0
191	840000.0	1093139.0	253139.0	30.0
688	865000.0	816184.0	-48816.0	6.0
10	930000.0	1435088.0	505088.0	54.0
420	1850000.0	834533.0	-1015467.0	55.0

```
[45]: sns.scatterplot(x=y_test, y=y_hat_test, alpha=0.6)  
sns.lineplot(y_test, y_test)  
  
plt.xlabel('Actual log wage', fontsize=14)  
plt.ylabel('Prediced log wage', fontsize=14)
```

```
plt.title('Actual vs Predicted log wage (test set)', fontsize=17)
plt.show()
```



```
[46]: sns.scatterplot(x=predictions['Predictions'], y=predictions['Residuals'],  
    ↪alpha=0.6)  
plt.show()  
  
# try to replicate this residual plot using log wages
```

2 Plotting the Simple Regression Line


Let's pick one of the most significant features and make a simple regression model with that.

```
[47]: df.head(4)
```


```
[47]:
```

	hours	IQ	educ	exper	tenure	age	married	black	meduc	lwage
0	40	93	12	11	2	31	1	0	8.0	13.552846
1	50	119	18	11	16	37	1	0	14.0	13.602317
2	40	108	14	11	9	33	1	0	14.0	13.623139
3	40	96	12	13	7	32	1	0	12.0	13.384728


```
[48]: sns.lmplot(x='educ',y='lwage',data=df)
plt.show()
```



```
[49]: sns.lmplot(x='educ',y='lwage',data=df, hue='married',palette='coolwarm') # how
      ↪do you interpret this one?
      plt.show()
```



```
[50]: sns.lmplot(x='educ',y='lwage',data=df, hue='black',palette='coolwarm') # how
      ↪do you interpret this one?
      plt.show()
```



```
[51]: sns.lmplot(x='educ',y='lwage',data=df, hue='black',palette='coolwarm',
               col='married') # how do you interpret this one?
plt.show()
```

